



Image Analysis and Object Recognition

Exercise 4

Summer Semester 2024

(Course materials for internal use only!)

Computer Vision in Engineering – Prof. Dr. Rodehorst

M.Sc. Mariya Kaisheva

mariya.kaisheva@uni-weimar.de

Agenda

Topics:

- Assignment 1.** Image enhancement, Binarization, Morphological operators
- Assignment 2.** Gradient of Gaussian filtering, Förstner interest operator
- Assignment 3.** Shape detection based on Hough-voting
- Assignment 4.** **Frequency domain filtering, Shape recognition via Fourier descriptors**
- Assignment 5.** Image segmentation using clustering
- Assignment 6.** Convolutional neural networks for image classification
- Final Project.** - *Will be announced during the last exercise class* -

Agenda

Start date and submission deadlines:

Assignment 1.	18.04.24 – 01.05.24
Assignment 2.	02.05.24 – 15.05.24
Assignment 3.	16.05.24 – 29.05.24
Assignment 4.	30.05.24 – 12.06.24
Assignment 5.	13.06.24 – 26.06.24
Assignment 6.	27.06.24 – 10.07.24
Final Project.	11.07.24 – 22.09.24

Wednesday by 23:00
(Central European Time)



Assignment 3: **Sample Solution**

Assignment 3: Overview

Topics:


- Hough line detection

Goal:

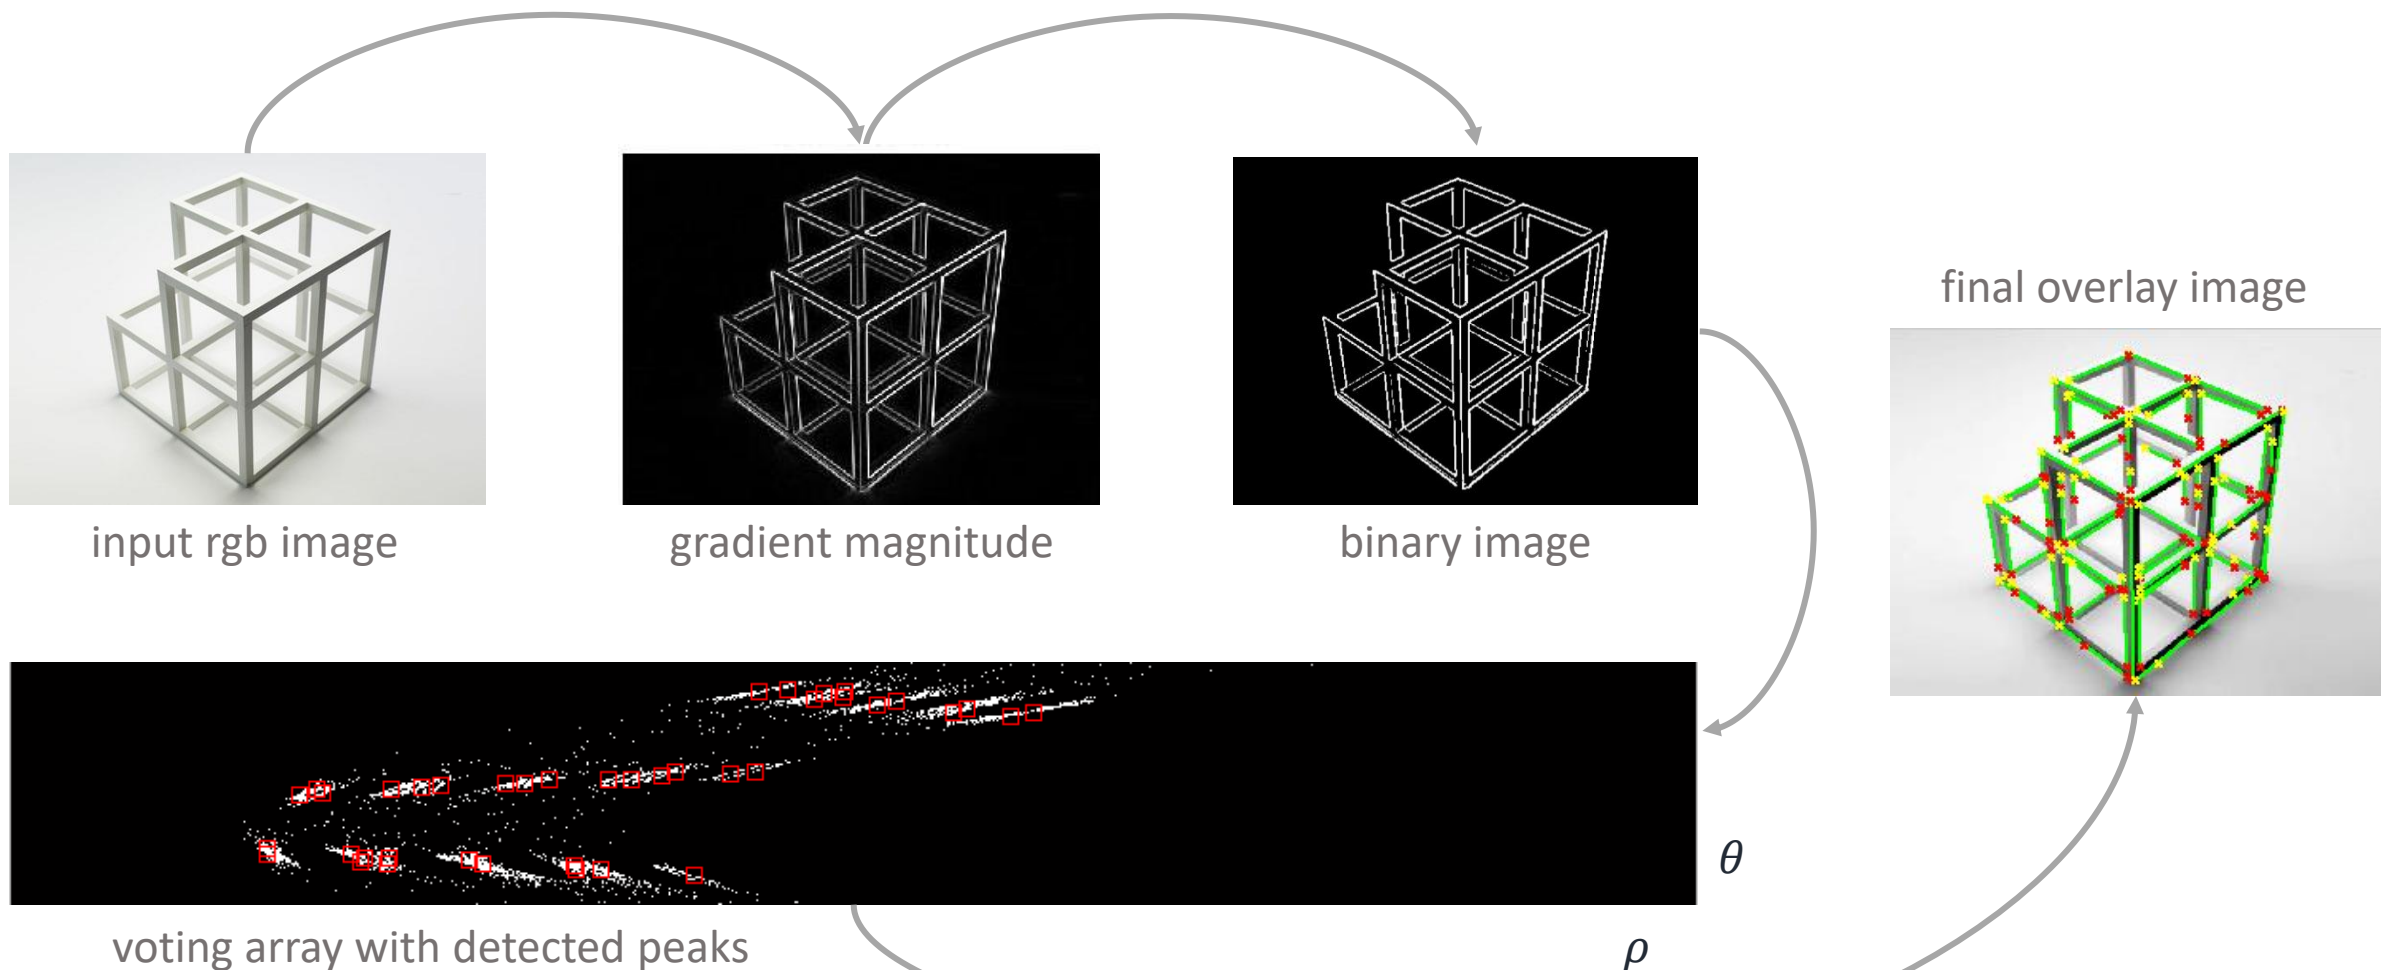
- Understanding the concept of Hough-voting
- Practice detection and parameterization of lines in images

Input:

- Provided image → *input_ex3.jpg*
- Or a different image of your own choice



Assignment 3: workflow



Algorithm outline

Input: binary edge image (from GoG-filtering)

Initialize index vectors

$$\rho_{ind} = [-\rho_{max}, \dots, \rho_{max}], \rho_{max} = \sqrt{n_{rows}^2 + n_{columns}^2}$$

$$\theta_{ind} = [-90, \dots, 89]$$

Initialize voting array H

$$H = \text{zeros}(2 \cdot \rho_{max} + 1, 180)$$

for each edge point (x, y) in the image

θ = gradient orientation at (x, y)


$$\rho = x \cdot \cos\theta + y \cdot \sin\theta$$

$$\theta_i = \text{find}(\theta_{ind} == \theta)$$

$$\rho_i = \text{find}(\rho_{ind} == \rho)$$

$$H(\rho_i, \theta_i) = H(\rho_i, \theta_i) + 1$$

end



main function

```
function Assignment3
% =====
sigma = 0.5; % standard deviation for smoothing
thres = 0.07; % binarization threshold
I = double(imread('input_ex3.jpg')) / 255;
figure; subplot(1, 4, 1); imshow(I); title('Original image');

[Ix, Iy] = Gradient(mean(I, 3), sigma); % image gradient from assignment 2
M = sqrt(Ix.^2 + Iy.^2); % gradient magnitude
subplot(1, 4, 2); imshow(M, []); title('Gradient magnitude');
BW = M > thres; % binary mask
subplot(1, 4, 3); imshow(BW); title('Binarized gradient');

[H, t, r] = my_hough(BW, Ix, Iy); % own Hough transform
subplot(1, 4, 4); imshow(imadjust(H), 'XData', t, 'YData', r);
title('Voting space'); ylabel('\rho'); xlabel('\theta');

peaks = houghpeaks(H, 40, 'threshold', 10); % find max values
hold on; plot(t(peaks(:,2)), r(peaks(:,1))), 's', 'color', 'red');

lines = houghlines(BW, t, r, peaks, 'FillGap', 5, 'MinLength', 10);
subplot(1, 4, 1); hold on;
for i = 1 : length(lines) % draw line segments
    xy = [lines(i).point1; lines(i).point2];
    plot(xy(:,1), xy(:,2), 'LineWidth', 2, 'Color', 'green');
    plot(xy(1,1), xy(1,2), 'x', 'LineWidth', 2, 'Color', 'yellow');
    plot(xy(2,1), xy(2,2), 'x', 'LineWidth', 2, 'Color', 'red');
end
```

`imshow(I, [])` displays the grayscale image `I`, scaling the display based on the range of pixel values in `I`.

`imshow(I)` displays the grayscale image `I`, using the default display range for the image data type.

helper function

```
function [H, t, r] = my_hough(BW, Ix, Iy)
% =====
d = round(sqrt(size(BW, 1)^2 + size(BW, 2)^2)); % image diagonal
t = -90:1:89; r = -d:1:d; % ranges for theta and rho
H = zeros(length(r), length(t)); % initialize Hough voting space
[y, x] = find(BW > 0); % relevant pixel positions
for i = 1 : length(x)
    theta = round(atan2(Iy(y(i), x(i)), Ix(y(i), x(i))) * 180/pi);
    rho = round(x(i) * cos(theta * pi/180) + y(i) * sin(theta * pi/180));
    ind_r = find(r == rho); % index lookup for rho
    ind_t = find(t == theta); % index lookup for theta
    H(ind_r, ind_t) = H(ind_r, ind_t) + 1; % vote for theta and rho
end
```

modified algorithm

$$\theta = \tan^{-1} \left(\frac{\frac{\partial f}{\partial y}}{\frac{\partial f}{\partial x}} \right)$$

atan2(y, x) computes
atan(y/x) for corresponding
elements of y and x

helper function

round up to the
next integer value

```
function [H, t, r] = my_hough(BW, Ix, Iy)
% =====
d = round(sqrt(size(BW, 1)^2 + size(BW, 2)^2));           % image diagonal
t = -90:1:89; r = -d:1:d;                                % ranges for theta and rho
H = zeros(length(r), length(t));                          % initialize Hough voting space
[y, x] = find(BW > 0);                                    % relevant pixel positions
for i = 1 : length(x)
    theta = round(atan2(Iy(y(i), x(i)), Ix(y(i), x(i))) * 180/pi);
    rho = round(x(i) * cos(theta * pi/180) + y(i) * sin(theta * pi/180));
    ind_r = find(r == rho);                                % index lookup for rho
    ind_t = find(t == theta);                              % index lookup for theta
    H(ind_r, ind_t) = H(ind_r, ind_t) + 1;                % vote for theta and rho
end
```

helper function

```
function [H, t, r] = my_hough(BW, Ix, Iy)
% =====
d = round(sqrt(size(BW, 1)^2 + size(BW, 2)^2));           % image diagonal
t = -90:1:89; r = -d:1:d;                                % ranges for theta and rho
H = zeros(length(r), length(t));                          % initialize Hough voting space
[y, x] = find(BW > 0);                                    % relevant pixel positions
for i = 1 : length(x)
    theta = round(atan2(Iy(y(i)), x(i)), Ix(y(i), x(i))) * 180/pi);
    rho = round(x(i) * cos(theta * pi/180) + y(i) * sin(theta * pi/180));
    ind_r = find(r == rho);                                % index lookup for rho
    ind_t = find(t == theta);                              % index lookup for theta
    H(ind_r, ind_t) = H(ind_r, ind_t) + 1;                % vote for theta and rho
end
```

degree → radian
conversion

cos() and *sin()*
require input in **radians**

alternative functions
which work with input
in **degrees** are
cosd() and *sind()*

main function

increases the
contrast →

```
function Assignment3
% =====
sigma = 0.5; % standard deviation for smoothing
thres = 0.07; % binarization threshold
I = double(imread('input_ex3.jpg')) / 255;
figure; subplot(1, 4, 1); imshow(I); title('Original image');

[Ix, Iy] = Gradient(mean(I, 3), sigma); % image gradient from assignment 2
M = sqrt(Ix.^2 + Iy.^2); % gradient magnitude
subplot(1, 4, 2); imshow(M, []); title('Gradient magnitude');
BW = M > thres; % binary mask
subplot(1, 4, 3); imshow(BW); title('Binarized gradient');

[H, t, r] = my_hough(BW, Ix, Iy); % own Hough transform
subplot(1, 4, 4); imshow(imadjust(H), 'XData', t, 'YData', r);
title('Voting space'); ylabel('\rho'); xlabel('\theta');

peaks = houghpeaks(H, 40, 'threshold', 10); % find max values
hold on; plot(t(peaks(:,2)), r(peaks(:,1))), 's', 'color', 'red');

lines = houghlines(BW, t, r, peaks, 'FillGap', 5, 'MinLength', 10);
subplot(1, 4, 1); hold on;
for i = 1 : length(lines) % draw line segments
    xy = [lines(i).point1; lines(i).point2];
    plot(xy(:,1), xy(:,2), 'LineWidth', 2, 'Color', 'green');
    plot(xy(1,1), xy(1,2), 'x', 'LineWidth', 2, 'Color', 'yellow');
    plot(xy(2,1), xy(2,2), 'x', 'LineWidth', 2, 'Color', 'red');
end
```



main function

```
function Assignment3
% =====
sigma = 0.5; % standard deviation for smoothing
thres = 0.07; % binarization threshold
I = double(imread('input_ex3.jpg')) / 255;
figure; subplot(1, 4, 1); imshow(I); title('Original image');

[Ix, Iy] = Gradient(mean(I, 3), sigma); % image gradient from assignment 2
M = sqrt(Ix.^2 + Iy.^2); % gradient magnitude
subplot(1, 4, 2); imshow(M, []); title('Gradient magnitude');
BW = M > thres; % binary mask
subplot(1, 4, 3); imshow(BW); title('Binarized gradient');

[H, t, r] = my_hough(BW, Ix, Iy); % own Hough transform
subplot(1, 4, 4); imshow(imadjust(H), 'XData', t, 'YData', r);
title('Voting space'); ylabel('\rho'); xlabel('\theta');

peaks = houghpeaks(H, 40, 'threshold', 10); % find max values
hold on; plot(t(peaks(:,2)), r(peaks(:,1)), 's', 'color', 'red');

lines = houghlines(BW, t, r, peaks, 'FillGap', 5, 'MinLength', 10);
subplot(1, 4, 1); hold on;
for i = 1 : length(lines) % draw line segments
    xy = [lines(i).point1; lines(i).point2];
    plot(xy(:,1), xy(:,2), 'LineWidth', 2, 'Color', 'green');
    plot(xy(1,1), xy(1,2), 'x', 'LineWidth', 2, 'Color', 'yellow');
    plot(xy(2,1), xy(2,2), 'x', 'LineWidth', 2, 'Color', 'red');
end
```

maximum
number of peaks

maximum value to
be considered a peak





Assignment 4

Assignment 4: Overview

Topics:

- Filtering in the frequency domain
- Shape recognition using Fourier descriptors

Goal:

- Practice noise removal in the frequency domain (Task A)
- Practice automatic shape detection using Fourier descriptors (Task B)


Input:

- All images provided for this assignment can be found on Moodle course page

Assignment 4: Image filtering in frequency domain

Task A: Image filtering

- Read the input image *taskA.png* and convert it to a grayscale image (double values between 0.0 and 1.0)
- Add Gaussian noise to the image (`imnoise`, parameters e.g. $M=0$, $V=0.01$) and plot the result
- Filter the noisy image with a self-made 2D Gaussian filter in the frequency-domain (`fft2`, `ifft2`). Which σ is suitable to remove the noise? Plot the result
- Plot the logarithmic centered image spectra of the noisy image, the (padded) Gaussian filter and the filtered image (`imagesc`, `log`, `abs` and `fftshift`)




Convert to grayscale


Add noise

Task A

$f(x, y)$




FFT
 \Rightarrow




$F(u, v)$

image spectrum




`imagesc(abs(fft_image))`

logarithmic scaled image spectrum



`imagesc(log(abs(fft_image)))`


centered scaled image spectrum



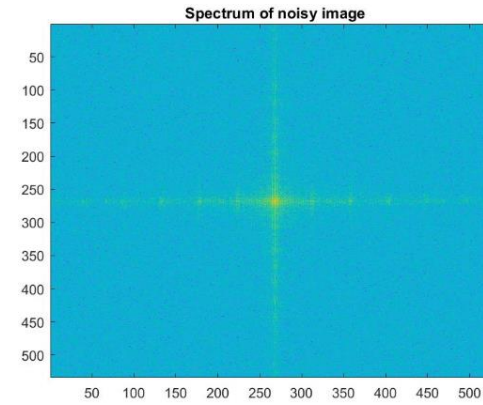
`imagesc(log(abs(fftshift(fft_image))))`

Task A

$$f(x, y)$$




FFT
 \Rightarrow




$$F(u, v)$$

Task A

$$f(x, y)$$




FFT
 \Rightarrow




$$F(u, v)$$

$$h(x, y)$$




Task A

$$f(x, y)$$




FFT
 \Rightarrow




$$F(u, v)$$

$$h(x, y)$$




Task A

$f(x, y)$




FFT
 \Rightarrow



$F(u, v)$


$h(x, y)$




Filter after
padding

Task A

$f(x, y)$




FFT
 \Rightarrow




$F(u, v)$

$h(x, y)$




Task A

$$f(x, y)$$




FFT
 \Rightarrow




$$F(u, v)$$

$$h(x, y)$$




FFT
 \Rightarrow



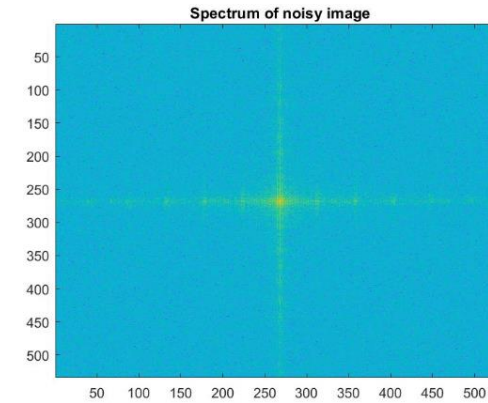
$$H(u, v)$$

Task A

$$f(x, y)$$




FFT
 \Rightarrow




$$F(u, v)$$

$$h(x, y)$$



FFT
 \Rightarrow

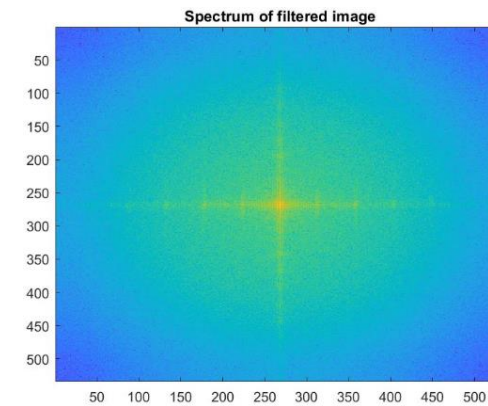


$$\cdot *$$

$$H(u, v)$$




$$G(u, v)$$




Task A

$$f(x, y)$$




FFT
 \Rightarrow

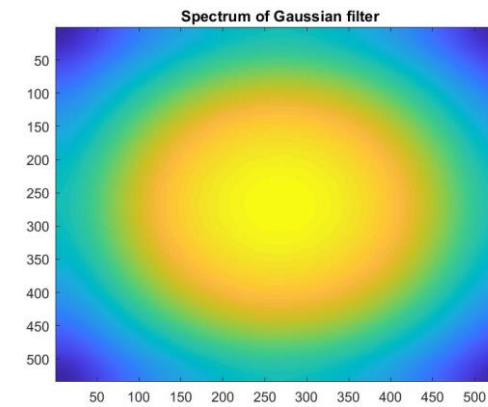


$$F(u, v)$$

$$h(x, y)$$



FFT
 \Rightarrow




$$H(u, v)$$

\cdot

$$g(x, y)$$



FFT⁻¹
 \Leftarrow



$$G(u, v)$$

\Downarrow

Assignment 4: Shape recognition

Task B: Image filtering

- Read the image `trainB.png` and convert it to a grayscale image (double values between 0.0 and 1.0)
- Derive a binary mask (data type `logical`) of the image where 1 represents the object of interest and 0 is background (`graythresh` and `im2bw`)
- Build a Fourier-descriptor D_f based on the binary mask of b.
 - Extraction of boundaries of the binary mask: `bwboundaries`
 - Use $n=24$ elements for the descriptor
 - Make it invariant against translation, orientation and scale
- Apply steps a.-c. on the images `test1B.jpg`, `test2B.jpg` and `test3B.jpg` in order to identify all potential object boundaries in the images. Note that here more than one boundaries will be identified by `bwboundaries`
- Identify the searched object by comparison of the trained Fourier-descriptor (result of task c) with all identified descriptors of the two test images (result of task d). Use the Euclidean distance of the Fourier-descriptors for identification, i.e.

$$\text{norm}(D_{f,\text{train}} - D_{f,\text{test}}) < 0.09$$

- Plot the identified boundaries on your mask (result of task b.) in order to validate the results (`imshow`, `hold on` and `plot`)

Task B

Input data



training image



test image 1




test image 2



test image 3

Boundary extraction

Task B



Fourier descriptor

Task B


- **Given:** m points representing the boundary of a **closed** region in the image
- Interpret the boundary coordinates (x, y) as complex numbers
- $b = \begin{bmatrix} (y_1, x_1) \\ \vdots \\ (y_m, x_m) \end{bmatrix}$ ($m \times 2$ array: output of *bwboundaries*)
- Build the **complex vector D** :

$$D = b(:, 2) + i * b(:, 1);$$

where $i^2 = -1$


Note: In MATLAB, the built-in **imaginary unit** can be found denoted both as i and as j

Don't use j / i as a variable in your code!




Reducing the number of elements n in $D_f \rightarrow$ shape generalization


Task B




$n = 2$



$n = 4$




$n = 8$



$n = 16$




$n = 24$




Reducing the number of elements n in $D_f \rightarrow$ shape generalization


Task B




$n = 2$



$n = 4$




$n = 8$



$n = 16$



$n = 24$



Extract only the **first n** elements (low frequency values) of the **Fourier-descriptor D_f** and ignore the rest

Implementation details

Task B

- c. Build a Fourier-descriptor based on the binary image of b.
 - i. Extraction of boundaries of the binary mask: `bwboundaries`
 - ii. Use $n = 24$ elements for the descriptor
 - iii. Make it invariant against translation, orientation and scale

Results:

- the final descriptor $D_{f,train}$ is a $1 \times n$ **vector**
- a 1×1 **cell** (matlab data type) containing an $m \times 2$ array which represent the m corresponding border pixel coordinates of the found shape (output of `bwboundaries`)

Task B

Output of bwboundaries:
($k \times 1$) cell,

where k is the number of
identified closed boundaries

```
My_Cell =  
  
[682x2 double]  
[686x2 double]  
[654x2 double]  
[685x2 double]  
[154x2 double]  
[168x2 double]  
[328x2 double]  
[335x2 double]  
[377x2 double]  
[332x2 double]  
[ 52x2 double]  
[333x2 double]  
[350x2 double]  
[288x2 double]  
[ 98x2 double]  
[196x2 double]  
[ 57x2 double]  
[ 41x2 double]  
[ 44x2 double]  
[189x2 double]  
[458x2 double]  
[326x2 double]  
[253x2 double]  
[ 84x2 double]  
[ 74x2 double]  
[244x2 double]  
[289x2 double]  
[209x2 double]  
[239x2 double]  
[ 87x2 double]  
[238x2 double]  
[ 84x2 double]  
[ 58x2 double]  
[ 12x2 double]  
[ 3x2 double]  
[216x2 double]
```

Access the 34th array of
boundary coordinates:

```
K>> boundary_points = My_Cell{34}  
  
boundary_points =  
  
    51    886  
    50    887  
    49    888  
    50    888  
    51    888  
    52    888  
    53    888  
    54    888  
    54    887  
    53    887  
    52    887  
    51    886
```

Implementation details

Task B

- d. Apply steps a.-c. on images *test1B.jpg*, *test2B.jpg* and *test3B.jpg* in order to identify all potential objects

Results for each image:

- **Descriptors:** $k \times n$ array, where k is the number of identified boundaries
- **Boundaries:** $k \times 1$ cell containing k ($m \times 2$) arrays which represent the corresponding border **pixel coordinates** of the k found shapes

Expected results

Task B



training image

